

Information Retrieval System

CSI 2290 Final Project

Aaron Kenward, Cameron Vogeli, Eric Daulo, Weiye Shi, Carlos Buenaventura, John Musa

Electrical and Computer Engineering Department

School of Engineering and Computer Science

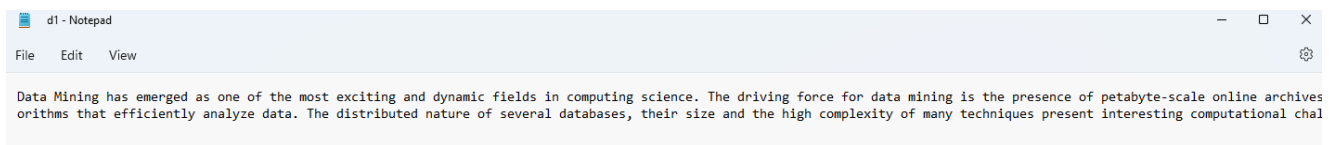
Oakland University, Rochester, MI

e-mails: akenward@oakland.edu, cmvogeli@oakland.edu, edaulo@oakland.edu , wshi@oakland.edu
caherrerabuenav@oakland.edu, johnmusa@oakland.edu

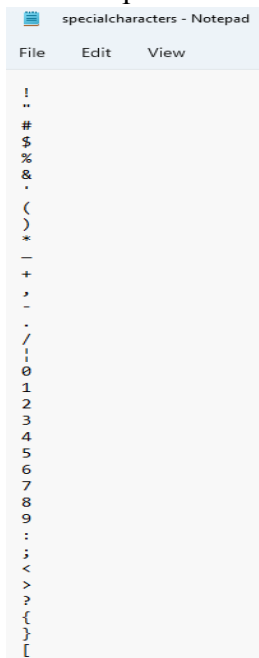
This project was designed to test students' abilities to read inputs from a file, parse the data, censor unwanted characters, and then tokenize them into a counted list file. This was to require skill in all disciplines of the course so far. After the dust had cleared with completing our code, the product was a complete tokenized file with words listed alphabetically and their instances counted. During this process, important functions like `qsort` and `strcmp` were discovered and relied upon to make quick work of the task. We concluded from this project that these kinds of algorithms are difficult and the people who write them every day in industry are extremely talented. The project was a mental, intellectual, and social exercise for all of us.

FILES WITH TEXT

As an example of the text files we were working with, this is the way `d1` appears:



This text file is full of characters, text or not, and we have to separate and count these. There is also files for special characters and stop words, like this:



READING AND WRITING DATA TO FILES

In the main function, we create a `FILE*` for each input file and use the `fopen` function to open the input files in read mode. We also tested if it was possible to open each file. If one of these files could not be opened, we write a “couldn’t open file” statement to `stderr` and exit the program with a 1. If they can be opened, we call our `init` function and pass each `FILE*` along with the path of the output file. At the end of our `init` function, we use `fclose` to close the `FILE*` passed to this function. The input and output files can be found under their own respective directories, all of which are under the `assets` folder.

To print each list to a file, we call the `printToFile` function called in `init` and pass the output file name and path as a `const char*`. In `printToFile`, we create a new `FILE*` and use `fopen` to open the output file in write mode. If the file cannot be opened we let the user know using `fprintf`. Otherwise, we iterate through the list of words and print each string to the file, separated by a new line.

When we first began the project, we didn’t create the `FILE*` in the main function. We passed the file names as `const char*` to each function in the project. From there, each function had to create a new `FILE*` that would be closed at the end of the block. For efficiency and do-it-once principles, we opted to instead create the `FILE*`’s once and keep them open until we are finished with them.

STRUCTURES

It was decided to use two main structure types for this project: `Str` and `Sqlist`. `Str` was used for storing the words themselves and `Sqlist` was used for tokenizing the words and storing list properties. `Str` possesses two pieces of data: a character array to store the word and a double value to store the number of times it appears. `Sqlist` is a bit more abstract and firstly stores the corresponding `Str` structure previously defined. It also contains value for the length which defines the size of the list for the words.

INITIALIZE LIST STRUCTURE

In order to utilize the list for counting and sorting the words, we first must initialize the values for this so it can run properly. This begins by creating a new structure of `Sqlist` type and passing it as a parameter to the `initSqlist` function. The other parameter is an integer value for the `maxsize`. For this function, the first step is to dynamically allocate memory for the list we will use. If the allocation is successful and passes our verification, then the `listsize` field of the new structure `L` is initialized to the `maxsize` value and the `length` field is set to 0.

PRINTING WORDS INTO THE LIST

After `initSqlist` is called, we then have to read the data from the input files into the list. To do this, we have a function called `getWords`. The function starts by creating character arrays (strings): one for storing the text from a file [called `string`], one as a buffer for storing the string before stop words and special characters are removed, and one for storing the stop words. The function runs through a loop using `fgets` until a null condition is reached (end of file).

In the loop, it first uses `strcpy()` to copy the contents of the buffer into `string`. This string is then delimited into words by using `strtok()` and setting spaces as a delimiter. A character-type pointer is then set equal to this delimited string. After this, a while loop is used to read through every line of the `stopwords.txt` file and it uses `strcmp` to see if any entries in the list are stop words. If they are not, then they get committed to `L` to be parsed and tokenized, which we will discuss later.

PARSING AND REMOVING SPECIAL CHARACTERS

A function called `parseWords` was used for this task. This function works by copying the data previously obtained through `getWords` into a buffer array. We then use a function called `strtok` which allows you to add delimiters as parameters, and we add all the special characters to that parameter. A loop is then used to remove the special characters systematically. It then adds the delimited word to the data field in `L` where it originally read from, and then iterates to the next data field to restart the cycle.

ALPHABETIZING THE LIST

Before we alphabetize the words in our list, we first needed to make the words lowercase. We achieved this by calling the `makeLowercase` function. In that function, we traverse through the entire list of words. In each word element, we iterate through its characters one by one, checking if the current character is uppercase. If it is, we first cast the character as an integer and add 32 to its value. Then, we cast it back into a char and assign it to the current character.

Now that all the words are lowercase, we alphabetize the list of words using `qsort`. We do so by passing to it the list of words, the length of the array, the size of the list of the words in bytes and the compare function. In the compare function, we use `strcmp` to compare two words and return a value back to `qsort`.

FINDING FREQUENCY AND WEIGHT

To determine frequency and weight of each word, we call `getFrequency` after we have gotten the words from the file, put them into an array, made each of them lowercase and alphabetized them. In the `getFrequency` function, we first find the frequency of each element in the list. We do this by creating two for loops and `strcmp`. Each indexed element is compared to the next element in the list using `strcmp`. If `strcmp` returns a 0, we write the total number of times each element occurs after this element. When the inner loop ends, we check if the number of times each element occurs is the greatest we found so far. If so, it becomes the new maximum frequency in the list.

After we are done iterating through the entire list, we go through it once more. We again use two loops, but this time if the current word is equal to the next word, we record the current word's frequency and skip to the next different word in the list. Before we do, we use `fprintf` to print the current word, its frequency and its weight.

The current code is the first edition of a function that calculates frequency and weight. Since the frequency is best found when the array is already sorted, this is one of the last functions we wrote.

CONCLUSIONS

After the conclusion of our work, many observations and areas for improvement could be discerned. We noticed that we had to find some external documentation to complete our work. This was a new process for some but the result benefited our work.

There is also the topic of room for improvement. When one runs the program, they can see 2 instances of special characters that slipped past our algorithm, so there can be refinement there. Some things could have also been made into bespoke functions instead of being called multiple times, which would've made the code more fail-proof. All in all, the code ran well and with a satisfactory result.

COMPILING AND RUNNING

To compile the project, navigate to the `src/` folder located in the project's root directory. In that folder, you should see a file called `main.c`. Compile this file using the normal means, naming the output file anything you wish. To run the file, execute the newly created binary. You should get an output similar to the following:

```
^  ~/De/csi2290-finalproject/src  on  root !4 ?1  gcc -o main main.c && ./main  at 17:28:57

=====
Current file: ../assets/outputs/Tokenizedd1.txt
=====
Word: $10          Frequency: 1.000000    Weight: 0.062500
Word: a            Frequency: 2.000000    Weight: 0.125000
Word: ad           Frequency: 1.000000    Weight: 0.062500
Word: algorithms   Frequency: 1.000000    Weight: 0.062500
Word: analysis     Frequency: 1.000000    Weight: 0.062500
Word: analyze      Frequency: 1.000000    Weight: 0.062500
Word: and          Frequency: 6.000000    Weight: 0.375000
Word: archives     Frequency: 1.000000    Weight: 0.062500
Word: around       Frequency: 1.000000    Weight: 0.062500
Word: artificial   Frequency: 1.000000    Weight: 0.062500
Word: as           Frequency: 1.000000    Weight: 0.062500
Word: assimilate   Frequency: 1.000000    Weight: 0.062500
Word: at           Frequency: 1.000000    Weight: 0.062500
Word: availability Frequency: 1.000000    Weight: 0.062500
Word: be           Frequency: 1.000000    Weight: 0.062500
Word: been         Frequency: 3.000000    Weight: 0.187500
Word: billion      Frequency: 1.000000    Weight: 0.062500
Word: bits         Frequency: 1.000000    Weight: 0.062500
Word: commercial   Frequency: 1.000000    Weight: 0.062500
Word: complexity   Frequency: 1.000000    Weight: 0.062500
Word: computational Frequency: 1.000000    Weight: 0.062500
Word: computing    Frequency: 1.000000    Weight: 0.062500
Word: concept      Frequency: 1.000000    Weight: 0.062500
Word: consequently Frequency: 1.000000    Weight: 0.062500
Word: contain      Frequency: 1.000000    Weight: 0.062500
Word: data         Frequency: 9.000000    Weight: 0.562500
Word: database     Frequency: 1.000000    Weight: 0.062500
Word: databases    Frequency: 2.000000    Weight: 0.125000
Word: detect       Frequency: 1.000000    Weight: 0.062500
Word: distributed  Frequency: 1.000000    Weight: 0.062500
Word: driving      Frequency: 1.000000    Weight: 0.062500
Word: dynamic      Frequency: 1.000000    Weight: 0.062500
Word: efficientl   Frequency: 1.000000    Weight: 0.062500
Word: emerged      Frequency: 1.000000    Weight: 0.062500
Word: enterprises  Frequency: 1.000000    Weight: 0.062500
Word: excess       Frequency: 1.000000    Weight: 0.062500
Word: exciting     Frequency: 1.000000    Weight: 0.062500
Word: expected     Frequency: 1.000000    Weight: 0.062500
Word: family       Frequency: 1.000000    Weight: 0.062500
Word: few          Frequency: 1.000000    Weight: 0.062500
Word: field        Frequency: 1.000000    Weight: 0.062500
Word: fields       Frequency: 1.000000    Weight: 0.062500
Word: for          Frequency: 3.000000    Weight: 0.187500
Word: force        Frequency: 1.000000    Weight: 0.062500
Word: form         Frequency: 1.000000    Weight: 0.062500
Word: has          Frequency: 1.000000    Weight: 0.062500
Word: have         Frequency: 3.000000    Weight: 0.187500
Word: hidden       Frequency: 1.000000    Weight: 0.062500
Word: high         Frequency: 1.000000    Weight: 0.062500
Word: hoc          Frequency: 1.000000    Weight: 0.062500
Word: in           Frequency: 5.000000    Weight: 0.312500
Word: information  Frequency: 1.000000    Weight: 0.062500
Word: intelligence Frequency: 1.000000    Weight: 0.062500
Word: interesting  Frequency: 2.000000    Weight: 0.125000
Word: intersection Frequency: 1.000000    Weight: 0.062500
Word: is           Frequency: 2.000000    Weight: 0.125000
```

Scrolling down, you should see each file's words along with their frequency and weights. Each input file is clearly marked as shown above. The output file of each input text file is located under

/assets/outputs/. Opening up “Tokenizedd1.txt” after executing the program should produce the following:

```
1  
2  
3 $10  
4 a  
5 a  
6 ad  
7 algorithms  
8 analysis  
9 analyze  
10 and  
11 and  
12 and  
13 and  
14 and  
15 and  
16 archives  
17 around  
18 artificial  
19 as  
20 assimilate
```

The other output files can be viewed in a similar manner.